

Amazon Web Services

The largest cloud provider, converting a decade of custom-silicon investment into an AI cost advantage · aws.amazon.com

2006

Founded

~28%

Global cloud share

\$37.6B

Q1 2026 revenue

\$25B

AI revenue run rate

\$364B

Contracted backlog

01 Overview

Holds about 28% of global cloud infrastructure against Azure's 21% and Google Cloud's 14%. The structural story is vertical integration: custom chips built over a decade now carry the majority of Bedrock inference, which converts directly into margin and price advantage as workloads shift from training to inference. Offering: Cloud infrastructure across 120 availability zones in 38 regions with 200+ services. The AI stack is Bedrock (managed multi-model inference), SageMaker (training), the Nova model family, and custom silicon — Trainium for AI, Graviton for CPU, Nitro for virtualisation.	Key facts HQ: Seattle, Washington (Amazon.com, Inc.) Ownership: Segment of Amazon.com, Inc. (NASDAQ: AMZN) Revenue: \$37.59B in Q1 2026 — roughly a \$150B annualised run rate — up from \$29.27B a year earlier, the fastest growth in 15 quarters. AI revenue was cited at a \$15B+ run rate in Q1 2026 and at \$25B by August 2026 Web: aws.amazon.com · LinkedIn Key people: Matt Garman (Chief Executive Officer, AWS) · Andy Jassy (President & Chief Executive Officer, Amazon (AWS founder)) Digital twin: Via partners (IoT TwinMaker) · GenAI: Yes (Bedrock, Nova)
---	--

02 Use Cases & Reported Metrics

Sector	Deployment	Reported impact	Source
AI platform	Amazon Bedrock, the managed multi-model inference service hosting Anthropic, Meta, Mistral, OpenAI and Amazon's own Nova models.	Processed more tokens in Q1 2026 than in all prior years combined; customer spend up 170% quarter over quarter	Link
Custom silicon	Trainium accelerators for training and inference, sold as capacity rather than chips, alongside continued NVIDIA GPU availability.	Powers over 50% of Bedrock token usage; 20–30% inference cost saving; Trainium2 ~30% better price-performance than comparable GPUs, Trainium3 a...	Link
Digital infrastructure	Data-centre build-out: \$15B committed to Northern Indiana facilities delivering 2.4 GW (December 2025) and a \$7B, 14-year framework agreement with Telangana, India to expand the Hyderabad footprint.	2026 capital expenditure guided to \$200B, later raised to \$220B	Link

03 Competitive Landscape

Other vendors in this market: Anthropic · Google · Meta · Microsoft · NVIDIA · OpenAI · Perplexity · xAI

Where it sits: AWS is the default substrate rather than a differentiated AI vendor — for an infrastructure portfolio the decision is rarely whether to use it but how much concentration risk to accept. The genuinely distinctive asset is Trainium: it powers more than 50% of Bedrock token usage, saves customers 20–30% on inference against market alternatives, and underpins a chip portfolio past a \$20B annual run rate.

04 GIP & BlackRock Connections

GIP became a wholly-owned BlackRock subsidiary in October 2024 and now operates as BlackRock's dedicated infrastructure-investment arm. Stakes below are standard institutional/index-fund positions unless otherwise noted.

Entity	Possible connection
Global Infrastructure Partners	No publicly disclosed fund-level relationship. AWS is near-certain to already be present inside multiple PortCos, so the practical question is consolidated commercial terms across the portfolio rather than a first introduction.
BlackRock	No publicly disclosed operating relationship; standard index holdings of AMZN only.

05 Sentiment

Customer: Demand is outrunning supply, which is the clearest possible demand signal and the worst possible negotiating environment. Garman has said Trainium is largely sold out through the end of next year, that AWS is securing five-year customer commitments, and that capacity is booked through 2028.

What supports the story	What to pressure-test
• Re-accelerating: Q1 2026 revenue up from \$29.27B to \$37.59B, the steepest growth in 15 quarters • Vertical integration is showing up in the P&L — custom silicon past a \$20B run rate, with management claiming several hundred basis points of margin advantage on inference versus third-party chips • \$364B contracted backlog, up 40% year over year, de-risks the capex programme	• \$200–220B of 2026 capital expenditure front-loads spend ahead of monetisation and compresses free cash flow • Capacity sold out through 2028 removes buyer leverage and slows onboarding for new workloads • Concentration risk cuts both ways — much of the AI revenue growth traces to a small number of frontier-lab counterparties

Employee: Not covered in this pass — AWS does not report headcount separately from Amazon, and employer-review data for Amazon as a whole is dominated by the fulfilment workforce and would not fairly represent AWS. Treat as a gap.

Sources: Amazon Web Services public filings/website, company fact sheets, Wikipedia, PitchBook/Crunchbase/Datanyze/BuiltIn profiles, Glassdoor employer reviews, vendor case studies and press releases as of Sept. 2026. Figures marked “est.” or “not disclosed” are third-party estimates or unavailable — not company-confirmed — validate in diligence. Part of the Physical AI / Industrial AI for Critical Infrastructure market landscape.