



Ali Ghodsi

Co-Founder & Chief Executive Officer

2013

Founded

\$7B

Run rate (Q2 2026)

>80%

YoY growth

\$190B

Valuation (Aug 2026)

\$1.5B

Lakehouse run rate

01 Who You're Meeting

Ali Ghodsi — Co-Founder & Chief Executive Officer, Databricks.

In role: CEO since January 2016; co-founded Databricks in 2013

Today: Runs Databricks' growth and international expansion, and is the public voice of the data-intelligence thesis — one governed source of enterprise truth with AI embedded throughout.

An academic computer scientist who turned out to be an exceptional operator. He holds a PhD in distributed computing from KTH Royal Institute of Technology in Sweden and an MBA, and came to UC Berkeley as a visiting scholar, joining the AMPLab group. He was one of the creators of Apache Spark, and ideas from his research on resource management, scheduling and data caching went into Apache Mesos and Apache Hadoop.

Career: Chief Executive Officer, Databricks; Co-Founder & VP of Engineering and Product Management, Databricks; Visiting scholar, AMPLab; co-creator of Apache Spark, UC Berkeley; Adjunct Professor; board member, RiseLab, UC Berkeley

Education: PhD Distributed Computing, KTH Royal Institute of Technology, Sweden (2006); MBA, Mid-Sweden University (2003)

The company: Grew from a \$5.4B run rate in February 2026 to \$7B by Q2, at over 80% growth, and raised \$5B at a \$190B valuation in August 2026 despite being cash-flow positive. The core data warehouse alone is \$1.5B of run rate still growing at 100%.

What they sell: The lakehouse platform combining data warehousing, engineering and ML, with Unity Catalog for governance. The newer agent layer is Lakebase (a database built for AI agents), Genie (a business-analysis agent), Agentbricks (embedding intelligence into software) and Unity AI Gateway (controlling model use and cost).

02 Where The Company Sits

Databricks' pitch to an infrastructure portfolio is neutrality and governance: it positions itself as model-agnostic, letting customers run OpenAI, Anthropic, Google or open-source models on their own data without that data leaving their secure cloud environment. For assets with sensitive operational data and multiple regulators, that architecture matters more than benchmark performance.

03 Talking Points

- A researcher-CEO: the governance and cost-control argument will get a more substantive hearing than a capability pitch.
- He is explicit that model choice is commoditising and the data layer is where the durable advantage sits — useful framing for a portfolio Data Foundation conversation.
- Databricks raised roughly \$20B in 20 months while cash-flow positive. Ask why; his board's answer was concern about an AI slump.

What to pressure-test about the company:

- Multibillion-dollar cloud commitments to all three hyperscalers create a fixed cost base regardless of demand
- Roughly \$20B raised over 20 months for a business that describes itself as cash-flow positive — the board's stated motivation was concern about a possible AI slump
- No S-1 filed and Ghodsi has said 2026 is a terrible year to go public; a listing is unlikely before Anthropic or OpenAI

04 GIP Context

GIP: No publicly disclosed GIP relationship. Worth noting that Blackstone participated in the August 2026 round — a peer infrastructure and private-markets investor taking a position in the vendor, which is context for any commercial discussion.

BlackRock: No publicly disclosed operating relationship. Databricks is private; the August 2026 round was led by Coatue with Blackstone, MGX, T. Rowe Price accounts and Sixth Street participating.

Others at Databricks: Matei Zaharia (Co-Founder & Chief Technology Officer)

www.databricks.com · [Databricks speaker profile](#) · [Databricks founders page](#) · [August 2026 funding round](#) · [Company LinkedIn](#)

Public professional and company information from the research layer of the GIP vendor tracker, generated September 15, 2026. Figures are as reported by the sources listed on the vendor's white page and are not independently audited. Research goes stale — re-run before external use.