



Matei Zaharia

Co-Founder & Chief Technology Officer

2013

Founded

\$7B

Run rate (Q2 2026)

>80%

YoY growth

\$190B

Valuation (Aug 2026)

\$1.5B

Lakehouse run rate

01 Who You're Meeting

Matei Zaharia — Co-Founder & Chief Technology Officer, Databricks.

In role: Co-founded Databricks in 2013 and served as its first CTO

Today: Owns the technology and research roadmap — Apache Spark, Delta Lake, MLflow, DBRX and the current work on combining large language models with external data sources.

He created Apache Spark as his PhD project at UC Berkeley's AMPLab in 2009, work that won the 2014 ACM Doctoral Dissertation Award and a US Presidential Early Career Award for Scientists and Engineers. He has since led or contributed to most of the open-source projects the modern data stack rests on — MLflow, Delta Lake and DBRX among them. He is an Associate Professor of Computer Science at UC Berkeley, having returned there after years at Stanford, and holds the CTO role at Databricks alongside it.

Career: Co-Founder & Chief Technology Officer, Databricks; Associate Professor of Computer Science, UC Berkeley; Assistant Professor of Computer Science, Stanford University; PhD researcher, AMPLab; created Apache Spark, UC Berkeley

Education: PhD Computer Science, UC Berkeley (2014 ACM Doctoral Dissertation Award)

The company: Grew from a \$5.4B run rate in February 2026 to \$7B by Q2, at over 80% growth, and raised \$5B at a \$190B valuation in August 2026 despite being cash-flow positive. The core data warehouse alone is \$1.5B of run rate still growing at 100%.

What they sell: The lakehouse platform combining data warehousing, engineering and ML, with Unity Catalog for governance. The newer agent layer is Lakebase (a database built for AI agents), Genie (a business-analysis agent), Agentbricks (embedding intelligence into software) and Unity AI Gateway (controlling model use and cost).

02 Where The Company Sits

Databricks' pitch to an infrastructure portfolio is neutrality and governance: it positions itself as model-agnostic, letting customers run OpenAI, Anthropic, Google or open-source models on their own data without that data leaving their secure cloud environment. For assets with sensitive operational data and multiple regulators, that architecture matters more than benchmark performance.

03 Talking Points

- The most technically credentialed person in this contact list; an architecture conversation will be substantive and worth preparing for.
- Still an active academic, so he engages with the limits of the technology as readily as its capabilities — useful for a realistic view of what agents can do on operational data.
- His research agenda is effectively Databricks' three-year roadmap. What he is publishing is what the product will do.

What to pressure-test about the company:

- Multibillion-dollar cloud commitments to all three hyperscalers create a fixed cost base regardless of demand
- Roughly \$20B raised over 20 months for a business that describes itself as cash-flow positive — the board's stated motivation was concern about a possible AI slump
- No S-1 filed and Ghodsi has said 2026 is a terrible year to go public; a listing is unlikely before Anthropic or OpenAI

04 GIP Context

GIP: No publicly disclosed GIP relationship. Worth noting that Blackstone participated in the August 2026 round — a peer infrastructure and private-markets investor taking a position in the vendor, which is context for any commercial discussion.

BlackRock: No publicly disclosed operating relationship. Databricks is private; the August 2026 round was led by Coatue with Blackstone, MGX, T. Rowe Price accounts and Sixth Street participating.

Others at Databricks: Ali Ghodsi (Co-Founder & Chief Executive Officer)

www.databricks.com · [Databricks author profile](#) · [Databricks founders page](#) · [Company LinkedIn](#)

Public professional and company information from the research layer of the GIP vendor tracker, generated September 15, 2026. Figures are as reported by the sources listed on the vendor's white page and are not independently audited. Research goes stale — re-run before external use.