



Aravind Srinivas

Co-Founder & Chief Executive Officer

2022

Founded

~1,400

Employees (est.)

~\$450M+

ARR

~\$20B

Valuation

780M

Monthly queries

01 Who You're Meeting

Aravind Srinivas — Co-Founder & Chief Executive Officer, Perplexity.

In role: Co-founded Perplexity in August 2022 and has led it since

Today: Leads Perplexity's product and strategy, including the citation-first answer engine, the Comet agentic browser, and the publisher revenue-share programme.

One of very few people to have done research at all three leading AI labs, which is the origin of Perplexity's model-agnostic design. He took a dual degree at IIT Madras and a PhD in computer science at UC Berkeley under Pieter Abbeel, publishing on contrastive learning, vision transformers, diffusion generative models and reinforcement learning. Along the way he interned at OpenAI in 2018 on policy gradient algorithms, at DeepMind in 2019 on large-scale contrastive learning, and at Google Research in 2020 and 2021 on vision models including Bottleneck Transformers and HaloNet.

Career: Co-Founder & Chief Executive Officer, Perplexity; Research Scientist — language models and diffusion generative models, OpenAI; Research intern — vision models, Google Research; Research intern — large-scale contrastive learning, DeepMind

Education: PhD Computer Science, University of California, Berkeley (advised by Pieter Abbeel); BTech and MS, Indian Institute of Technology Madras

The company: Went from a \$121M valuation in April 2023 to roughly \$20B by September 2025, on annualised revenue reported above \$450M by March 2026. The strategic bet is Comet — a free agentic browser — on the thesis that owning the browsing surface beats owning the underlying model.

What they sell: An answer engine returning synthesised, cited answers rather than links, sold through a free tier, Pro at \$20/month, Max at \$200/month, enterprise plans, an API, and a publisher revenue-share programme. Comet, a Chromium-based agentic browser, went free worldwide in October 2025. Deep Research provides asynchronous multi-step research.

02 Where The Company Sits

Perplexity does not train frontier models; it buys inference from Anthropic, OpenAI and open-weight providers and competes on interface, citation discipline and distribution. That is a real product advantage for research-shaped work — compare X to Y, what does this paper say — and a structural vulnerability, because its suppliers are also its competitors and its valuation is priced where foundation labs trade.

03 Talking Points

- A researcher-founder who buys inference rather than training frontier models — the model-agnostic argument is his genuine position, not a hedge.
- Citation-first matters for any use where output has to be auditable. That is the case worth testing for research-heavy roles at PortCos.
- Publisher litigation is unresolved and material to enterprise adoption. He launched the revenue-share programme in response; ask where that now stands.

What to pressure-test about the company:

- Buys most frontier inference from Anthropic, OpenAI and open-weight providers — its suppliers are also its competitors
- Valued at roughly 40 to 100 times ARR depending on which figures are used, priced where foundation labs trade despite not training frontier models
- Active publisher lawsuits over content scraping are unresolved and material to enterprise adoption

04 GIP Context

GIP: No publicly disclosed GIP relationship. Relevance sits in the Individual Productivity area of the reference model — research and analysis workflows — rather than asset optimisation. Note that active publisher litigation over content scraping is a live consideration for any enterprise deployment.

BlackRock: No publicly disclosed operating relationship.

Others at Perplexity: Denis Yarats (Co-Founder & Chief Technology Officer)

www.perplexity.ai · [Forbes profile](#) · [TechCrunch author profile](#) · [Career analysis](#) · [Company LinkedIn](#)

Public professional and company information from the research layer of the GIP vendor tracker, generated September 15, 2026. Figures are as reported by the sources listed on the vendor's white page and are not independently audited. Research goes stale — re-run before external use.